


# *Lect 7* Text Classification


# Why do we need to classify texts?

- As a self-sufficient task:
  - Spam filtering
  - Sentiment analysis



**Is Twitter better at predicting elections than opinion polls?**

# Text classification in general



# Text label kinds

1. Discrete labels:
  - a. Label is known:
    - i. Binary classification: spam filtering/sentiment analysis
    - ii. Multi-class classification: categorization of goods
    - iii. Multi-label classification: #hashtag prediction
  - b. Label is unknown:
    - i. Text clasterization: user intent search
2. Continuous labels: predict a salary by CV, predict a price by a product description

# Classifier - Naïve Bayes (1)

- Given document  $d$  and a fixed set of classes  $\mathcal{C} = \{c_1, c_2\}$ ,  $x$  is the words of  $d$ , compute the class of  $d$ :

$$c_{MAP} = \operatorname{argmax}_{c \in \mathcal{C}} P(c|d)$$

$$\begin{aligned} &= \operatorname{argmax}_{c \in \mathcal{C}} \frac{P(d|c)P(c)}{P(d)} \\ &= \operatorname{argmax}_{c \in \mathcal{C}} P(d|c)P(c) \end{aligned}$$

Bayes rule

$$= \operatorname{argmax}_{c \in \mathcal{C}} P(x_1, x_2, \dots, x_n | c)P(c)$$

$$= \operatorname{argmax}_{c \in \mathcal{C}} P(c) \prod_{x \in X} P(x|c)$$

# Classifier – Naïve Bayes (2)

$$\hat{P}(c) = \frac{N_c}{N}$$

$$\hat{P}(w|c) = \frac{\text{count}(w,c)+1}{\text{count}(c)+|V|}$$

	Doc	Words	Class
Training	1	Chinese Beijing Chinese	c
	2	Chinese Chinese Shanghai	c
	3	Chinese Macao	c
	4	Tokyo Japan Chinese	j
Test	5	Chinese Chinese Chinese Tokyo Japan	?

**Conditional Probabilities:**

**Priors:**

$$P(c) = \frac{3}{4} \quad P(j) = \frac{1}{4}$$

$$P(\text{Chinese}|c) = (5+1) / (8+6) = 6/14 = 3/7$$

$$P(\text{Tokyo}|c) = (0+1) / (8+6) = 1/14$$

$$P(\text{Japan}|c) = (0+1) / (8+6) = 1/14$$

$$P(\text{Chinese}|j) = (1+1) / (3+6) = 2/9$$

$$P(\text{Tokyo}|j) = (1+1) / (3+6) = 2/9$$

$$P(\text{Japan}|j) = (1+1) / (3+6) = 2/9$$

**Choosing a class:**

$$P(c|d5) \propto 3/4 * (3/7)^3 * 1/14 * 1/14$$

$$\approx 0.0003$$

$$P(j|d5) \propto 1/4 * (2/9)^3 * 2/9 * 2/9$$

$$\approx 0.0001$$


from Dan Jurafsky and Christopher Manning, Stanford University

## Use cases for feedforward networks

Let's consider 2 (simplified) sample tasks:

1. Text classification ↗
2. Language modeling

# Neural Net Classification with embeddings as input features!



# Simple feedforward Neural Language Models

**Task:** predict next word  $w_t$


given prior words  $w_{t-1}, w_{t-2}, w_{t-3}, \dots$

**Problem:** Now we're dealing with sequences of arbitrary length.



**Solution:** Sliding windows (of fixed length)

$$P(w_t | w_1^{t-1}) \approx P(w_t | w_{t-N+1}^{t-1})$$

# Neural Language Model




# 1 Dimensional CNN | Conv1D (e.g. Text Data)



- 1D CNNs are also applied on:
  - Sensory data, (e.g. accelerometer data)
  - Audio Data
  - Text data

Since we can also represent the Sensor Readings, Sound and Texts as a time series data.


# 1 Dimensional CNN | Conv1D (e.g. Text Data)



# CNNs for Text Classification

## [3] Word Embedding

- Word embeddings are vectors of a specified length, typically on the order of 100s
- Each vector of 100 or so values, represents one word.
- The values in each column represent the features of a word, rather than any specific word.
- These embeddings are formed in an unsupervised manner by training a neural network (a Word2Vec model) on an input word and a few surrounding words in a sentence.




# CNNs for Text Classification

## [4] Convolutional Kernel

- Convolution in TEXT
- Convolutional kernel will still be a sliding window, only its job is to look at embeddings for multiple words, rather than small areas of pixels in an image.
- To look at sequences of word embeddings, we want a window to look at multiple word embeddings in a sequence.
- The kernels will no longer be square, instead, they will be a wide rectangle with dimensions like 3x300 or 5x300 (assuming an embedding length of 300).
- The **height of the kernel** will be the number of embeddings it will see at once, similar to representing an n-gram in a word model.
- The **width of the kernel** should span the **length** of an entire **word embedding**.





## Recognizing General Patterns


- Recall that similar words will have similar embeddings and a convolution operation is just a linear operation on these vectors.
- when a convolutional kernel is applied to different sets of similar words, it will produce a similar output value!
- the convolutional output value for the input 2-grams “**good movie**” and “**fantastic song**” are about the same because the word embeddings for those pairs of words are also very similar.



# CNNs for Text Classification

## 1-D Convolution





# Example

We build a CNN model that converts words into vectors, selects important features using pooling and combines them in fully connected layers. Dropout prevents overfitting and the final layer outputs a probability for classification.

**models.Sequential()**: Creates a linear stack of layers where each layer passes output to the next.

**layers.Embedding(input\_dim=10000, output\_dim=100, input\_length=500)**: Converts word indices into 100-dimensional vectors, helping the model learn word meanings. Handles a vocabulary of 10,000 words and sequences of 500 words.

**layers.Conv1D(filters=128, kernel\_size=5, activation='relu')**: Applies 128 sliding filters that look at 5 words at a time to detect patterns.

**layers.GlobalMaxPooling1D()**: Reduces data by taking the maximum value from each filter's output, keeping only the most important features.

**layers.Dense(64, activation='relu')**: A fully connected layer with 64 neurons that learns complex patterns.

**layers.Dropout(0.5)**: Randomly disables 50% of neurons during training to prevent overfitting.

**layers.Dense(1, activation='sigmoid')**: Final output layer that predicts a probability (0–1) for binary classification.

# Importing Libraries .1

We will import the required libraries such as tensorflow, numpy required for building CNN model , creating layers, handling numerical operations and padding text sequences.

- tensorflow.keras:** Used for importing layers like Embedding, Conv1D and Sequential for model building.
- imdb:** Loads the IMDB dataset.
- pad\_sequences:** Pads text sequences to a fixed length.

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, Conv1D, GlobalMaxPooling1D, Dense, Dropout
from tensorflow.keras.datasets import imdb
from tensorflow.keras.preprocessing import sequence
```

## Loading Data .2

We will load and preprocess the IMDB dataset.

- imdb.load\_data(num\_words=10000)**: Loads the IMDB dataset, keeping only the 10,000 most frequent words.
- pad\_sequences(sequences, maxlen=500)**: Pads or cuts reviews so each is exactly 500 words long.

```
vocab_size = 10000
```

```
max_length = 500
```

```
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=vocab_size)
```

```
x_train = sequence.pad_sequences(x_train, maxlen=max_length)
```

```
x_test = sequence.pad_sequences(x_test, maxlen=max_length)
```

# Building CNN model .3

We build a CNN model that converts words into vectors, selects important features using pooling and combines them in fully connected layers. Dropout prevents overfitting and the final layer outputs a probability for classification.

- models.Sequential()**: Creates a linear stack of layers where each layer passes output to the next.
- layers.Embedding(input\_dim=10000, output\_dim=100, input\_length=500)**: Converts word indices into 100-dimensional vectors, helping the model learn word meanings. Handles a vocabulary of 10,000 words and sequences of 500 words.
- layers.Conv1D(filters=128, kernel\_size=5, activation='relu')**: Applies 128 sliding filters that look at 5 words at a time to detect patterns.
- layers.GlobalMaxPooling1D()**: Reduces data by taking the maximum value from each filter's output, keeping only the most important features.
- layers.Dense(64, activation='relu')**: A fully connected layer with 64 neurons that learns complex patterns.
- layers.Dropout(0.5)**: Randomly disables 50% of neurons during training to prevent overfitting.
- layers.Dense(1, activation='sigmoid')**: Final output layer that predicts a probability (0–1) for binary classification.

```
model = Sequential([
    Embedding(vocab_size, 100, input_length=max_length),
    Conv1D(filters=128, kernel_size=5, activation='relu'),
    GlobalMaxPooling1D(),
    Dense(64, activation='relu'),
    Dropout(0.5),
    Dense(1, activation='sigmoid')
])
```

## .4 Compiling and Training the Model

We will compile the model and train it using the IMDB dataset.

Here we will use [Adam](#) optimizer with [binary cross-entropy](#) as loss function.

- **model.compile(optimizer='adam', loss='binary\_crossentropy', metrics=['accuracy'])**: Defines the optimizer (Adam), loss function (binary cross-entropy) and accuracy metric for evaluating performance.
- **model.fit(x\_train, y\_train, epochs=5, batch\_size=128, validation\_split=0.2)**: Trains the model for 5 epochs using batches of 128 samples, with 20% of the training data reserved for validation.

```
model.compile(optimizer='adam',
  loss='binary_crossentropy',
  metrics=['accuracy'])
model.fit(x_train, y_train, batch_size=32,
  epochs=5, validation_split=0.2)
```

## 5. Evaluating the Model

We will evaluate the trained model on the test dataset.

**model.evaluate(x\_test, y\_test)**: Evaluates model performance by returning loss and accuracy.

**print(f"Test Accuracy: {test\_accuracy:.4f}")**: Prints the accuracy percentage on the test data.

```
test_loss, test_accuracy = model.evaluate(x_test, y_test)
```

```
print(f"Test Accuracy: {test_accuracy:.4f}")
```

```
Epoch 1/5
625/625 111s 173ms/step - accuracy: 0.6519 - loss: 0.5842 - val_accuracy: 0.8842 - val_loss: 0.2808
Epoch 2/5
625/625 140s 171ms/step - accuracy: 0.9219 - loss: 0.2098 - val_accuracy: 0.9020 - val_loss: 0.2517
Epoch 3/5
625/625 143s 172ms/step - accuracy: 0.9773 - loss: 0.0777 - val_accuracy: 0.9006 - val_loss: 0.2752
Epoch 4/5
625/625 142s 171ms/step - accuracy: 0.9945 - loss: 0.0248 - val_accuracy: 0.8924 - val_loss: 0.3894
Epoch 5/5
625/625 139s 167ms/step - accuracy: 0.9986 - loss: 0.0081 - val_accuracy: 0.8988 - val_loss: 0.4441
782/782 35s 45ms/step - accuracy: 0.8860 - loss: 0.4716
Test Accuracy: 0.8884
```